

Digital System Design Lab Report, Lab6

Author: Ying Yiwen

Number: 12210159

Abstract

This lab report investigates three distinct multiplier architectures—combinational, repetitive-addition, and pipelined designs—for FPGA-based digital systems. The study aims to evaluate their trade-offs in speed, hardware resource utilization, and throughput. Combinational logic achieves maximum speed through parallel partial product generation but incurs exponential area complexity. Repetitive-addition minimizes resource usage via sequential addition and shift operations, albeit with slower computation. Pipelined design balances speed and throughput by segmenting operations into clock-synchronized stages, enabling parallel processing. Experimental results quantify these trade-offs: combinational design exhibits the fastest single-operation latency, pipelined design maximizes throughput, and repetitive-addition optimizes space efficiency. This analysis provides empirical insights for selecting multiplier architectures based on application-specific constraints in reconfigurable computing systems.

Contents

1	Introduction	2
2	Combinational Design	2
3	Repetitive-Addition Design	4
4	Pipelined Design	7
5	Code Appendix	10
6	Analysis	11
6.1	Speed	11
6.2	Space Usage	11
6.3	Throughput	12
7	Conclusion	12

1 Introduction

Digital arithmetic circuits form the cornerstone of modern FPGA-based computing systems, with multiplier units being particularly critical in applications ranging from signal processing to machine learning accelerators. This experiment explores three distinct multiplier implementations in VHDL: combinational logic, repetitive-addition, and pipelined architectures, each embodying fundamental trade-offs between temporal efficiency, hardware resource utilization, and power consumption.

Theoretical analysis suggests that combinational multipliers achieve maximal throughput through parallel partial product generation and carry propagation, albeit at the cost of exponential area complexity $O(n)^2$. In contrast, repetitive-addition designs leverage sequential finite state machines to iteratively accumulate results, sacrificing speed for linear area scaling $O(n)$. Pipelined architectures introduce temporal segmentation through register staging, theoretically enabling sub-linear latency amortization while maintaining moderate resource overhead.

This systematic comparison aims to quantify how these architectural paradigms manifest in physical FPGA implementations. Key metrics include LUT (Look-Up Table) consumption, maximum clock frequency, and pipeline initiation interval – parameters directly influencing application-specific design choices. Through this experimental framework, we establish empirical correlations between algorithmic complexity, hardware pragmatism, and performance scalability in reconfigurable computing environments.

2 Combinational Design

The combinational design is based on the principles of Bool algebra and combinational logic circuits. Multiplication operations can be realized directly through logic gates. Multiplication can be decomposed into generation and accumulation of partial products, and these operations can be represented by Bool expressions and realized by combinational logic circuits.

$$P = A \times B = \sum_{i=0}^{n-1} (A_i \times B) \times 2^i \quad (1)$$

where A_i is the i -th bit of A , B is the multiplier, and 2^i represents the shift operation.

The main goal of the combinational design is to achieve the fastest multiplication operation. By reducing the latency of logic gates and optimizing the circuit structure, the operation time can be minimized. However, this design usually sacrifices hardware resources because a large number of logic gates are required to implement complex multiplication logic.

The following is the VHDL code for the combinational design:

```
library IEEE;
use IEEE.STD_LOGIC_1164.ALL;
use IEEE.STD_LOGIC_ARITH.ALL;
use IEEE.STD_LOGIC_UNSIGNED.ALL;

entity combinational_design is
    port (
        a, b : in std_logic_vector(4 downto 0); -- Input vectors a and b, both 5 bits
            wide
        y : out std_logic_vector(9 downto 0)      -- Output vector y, 10 bits wide to hold
            the product
    );
end combinational_design;
```

```

architecture Behavioral of combinational_design is

    constant WIDTH : integer := 5; -- Define the width of the input vectors
    signal au, bv0, bv1, bv2, bv3, bv4: std_logic_vector (WIDTH-1 downto 0); --
        Intermediate signals
    signal pp0, pp1, pp2, pp3, pp4: std_logic_vector (WIDTH downto 0); -- Signals for
        partial products
    signal prod: std_logic_vector (2*WIDTH-1 downto 0); -- Signal to hold the final
        product

begin

    au <= a; -- Assign input a to signal au

    -- Create vectors bv0 to bv4 where each bit is replicated from the corresponding
        bit of b
    bv0 <= (others => b(0));
    bv1 <= (others => b(1));
    bv2 <= (others => b(2));
    bv3 <= (others => b(3));
    bv4 <= (others => b(4));

    -- Generate the first partial product pp0, Perform bitwise AND between bv0 and au,
        and prepend a '0'
    pp0 <= '0' & (bv0 and au);

    -- Generate the second partial product pp1, Shift pp0 right by one bit and add it
        to the result of bv1 AND au
    pp1 <= ('0' & pp0(WIDTH downto 1)) + ('0' & (bv1 and au));

    -- Generate the third partial product pp2, Shift pp1 right by one bit and add it
        to the result of bv2 AND au
    pp2 <= ('0' & pp1(WIDTH downto 1)) + ('0' & (bv2 and au));

    -- Generate the fourth partial product pp3, Shift pp2 right by one bit and add it
        to the result of bv3 AND au
    pp3 <= ('0' & pp2(WIDTH downto 1)) + ('0' & (bv3 and au));

    -- Generate the fifth partial product pp4, Shift pp3 right by one bit and add it
        to the result of bv4 AND au
    pp4 <= ('0' & pp3(WIDTH downto 1)) + ('0' & (bv4 and au));

    -- Concatenate the least significant bits of all partial products to form the
        final product
    prod <= pp4 & pp3(0) & pp2(0) & pp1(0) & pp0(0);

    -- Assign the final product to the output y
    y <= prod;

end Behavioral;

```

After synthesis and implementations, we can get the schematic diagram as follows:

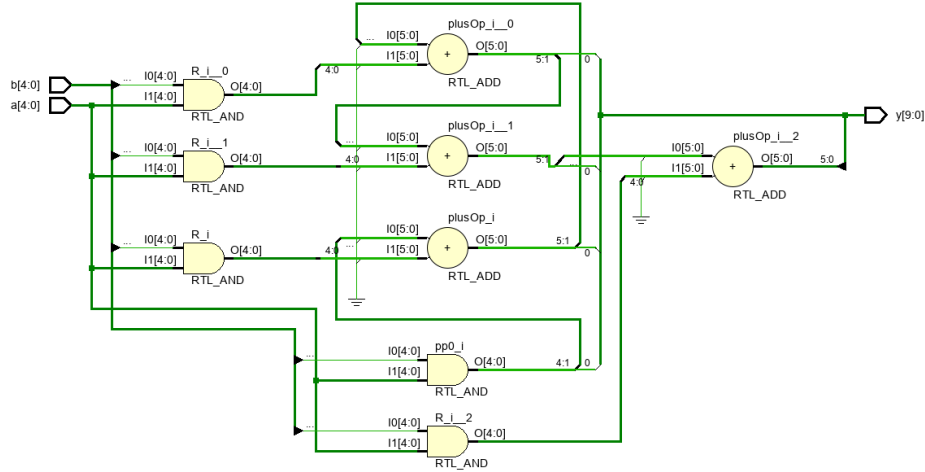


Figure 1: Combinational Design Schematic

Using the testbench in the appendix, we can get the result. It computes rightly.




Figure 2: Combinational Design Result

3 Repetitive-Addition Design

Repeated addition design is based on basic addition and shift operations. Multiplication can be viewed as a repeated execution of addition. $A \times B$ can be represented as adding A to B times. The weight assignment in multiplication can be realized by shift operation.

The main goal of a repetitive-addition design is to reduce the use of hardware resources. By using simple adders and shift registers, a smaller chip area and lower power consumption can be achieved, but at the expense of computing speed.

The following is the VHDL code for the repetitive-addition design:

```
library IEEE;
use IEEE.STD_LOGIC_1164.ALL;
use IEEE.STD_LOGIC_ARITH.ALL;
use IEEE.STD_LOGIC_UNSIGNED.ALL;

-- Entity declaration
entity repetitive_addition is
    port (
        CLK      : in  std_logic;           -- System clock
        RESET    : in  std_logic;           -- Active-high reset
        start    : in  std_logic;           -- Start computation signal
        a_in     : in  std_logic_vector(4 downto 0); -- Multiplicand input
        b_in     : in  std_logic_vector(4 downto 0); -- Multiplier input
        r        : out std_logic_vector(9 downto 0); -- 16-bit result output
        ready    : out std_logic            -- Computation complete flag
    );
end entity;
```

```

);
end entity repetitive_addition;

architecture Behavioral of repetitive_addition is
    constant WIDTH : integer := 5;           -- Data width configuration
    type state_type is (idle, ab0, load, op); -- FSM states:
        -- idle: Initial/waiting state
        -- ab0:  Handle zero input case
        -- load: Initialize registers
        -- op:   Multiplication operation state

    -- State and data registers
    signal state_reg, state_next : state_type;
    signal a_reg, a_next         : std_logic_vector(WIDTH-1 downto 0); -- Multiplicand
                                storage
    signal n_reg, n_next         : std_logic_vector(WIDTH-1 downto 0); -- Counter
                                storage
    signal r_reg, r_next         : std_logic_vector(2*WIDTH-1 downto 0); -- Result
                                accumulation

begin
    -- Synchronous State and Data Registers Update Process
    process (CLK, RESET) is
    begin
        if RESET = '1' then                -- Asynchronous reset
            state_reg <= idle;              -- Return to initial state
            a_reg      <= (others => '0');   -- Clear multiplicand register
            n_reg      <= (others => '0');   -- Clear counter register
            r_reg      <= (others => '0');   -- Clear result register
        elsif CLK'event and CLK = '1' then -- Clock edge triggered update
            state_reg <= state_next;        -- Update state register
            a_reg      <= a_next;            -- Update multiplicand register
            n_reg      <= n_next;            -- Update counter register
            r_reg      <= r_next;            -- Update result register
        end if;
    end process;

    -- Combinational Next-State and Data Path Logic
    process (start, state_reg, a_reg, n_reg, r_reg, a_in, b_in) is
    begin
        -- Default register values (prevent latches)
        a_next <= a_reg;
        n_next <= n_reg;
        r_next <= r_reg;
        ready  <= '0'; -- Default ready signal state

        case state_reg is
            when idle =>
                ready <= '1'; -- System ready for new operation
                if start = '1' then
                    -- Handle zero input special case
                    if (a_in = x"00" or b_in = x"00") then
                        state_next <= ab0;
                    else

```

```

        state_next <= load;  -- Proceed to normal operation
    end if;
else
    state_next <= idle;  -- Maintain idle state
end if;

when ab0 =>
    -- Immediate result for zero input case
    a_next <= a_in;      -- Store multiplicand
    n_next <= b_in;      -- Store multiplier
    r_next <= (others => '0');  -- Clear result
    state_next <= idle;  -- Return to idle

when load =>
    -- Initialize registers for multiplication
    a_next <= a_in;      -- Load multiplicand
    n_next <= b_in;      -- Initialize counter
    r_next <= (others => '0');  -- Reset accumulator
    state_next <= op;    -- Proceed to operation

when op =>
    -- Core multiplication iteration
    n_next <= n_reg - 1;  -- Decrement counter
    r_next <= ("0000" & a_reg) + r_reg;  -- Accumulate sum

    -- Check completion condition
    if (n_reg = x"01") then  -- Last iteration
        state_next <= idle;  -- Return to idle
    else
        state_next <= op;    -- Continue iterations
    end if;
end case;
end process;

-- Connect internal register to output port
r <= r_reg;

end architecture Behavioral;

```

After synthesis and implementations, we can get the schematic diagram as follows:

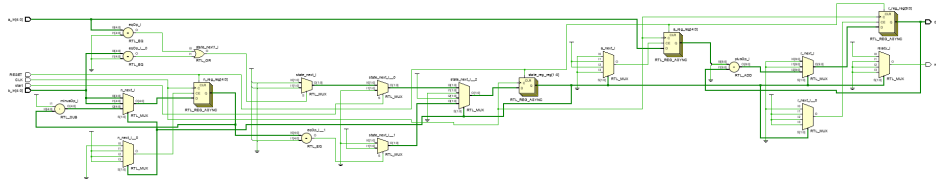


Figure 3: Repetitive-Addition Design Schematic

Using the testbench in the appendix, we can get the result. It computes rightly.

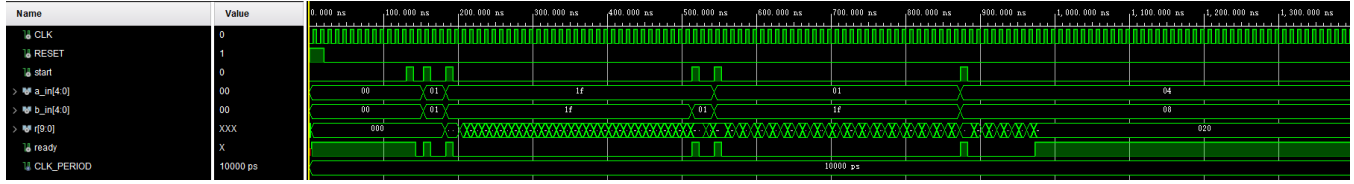


Figure 4: Repetitive-Addition Design Result

4 Pipelined Design

The pipeline design is based on the pipeline principle, which decomposes a complex multiplication operation into several simple stages and inserts registers between each stage. Each stage can be completed in a single clock cycle, thus enabling parallel processing.

The main goal of the pipelined design is to increase the throughput rate, i.e., the number of multiplication operations completed per unit of time. With parallel processing, the pipelined design can significantly improve the performance of the system by outputting a multiplication result in each clock cycle.

For an n -bit multiplier, if it's divided into m pipeline stages, each with a delay of T_{stage} , the delay of a single operation is:

$$T_{total} = T_{stage} \times m \quad (2)$$

The throughput rate of a pipeline design can be expressed as:

$$\text{Throughput} = \frac{1}{T_{clock}} \quad (3)$$

The following is the VHDL code for the pipelined design:

```
library IEEE;
use IEEE.STD_LOGIC_1164.ALL;
use IEEE.STD_LOGIC_ARITH.ALL;      -- Legacy arithmetic package
use IEEE.STD_LOGIC_UNSIGNED.ALL;   -- Unsigned operations

-- Entity Declaration
entity pipeline_multiplier is
  generic(WIDTH : integer := 5);    -- Default 5-bit operand width
  port(
    clk      : in  std_logic;        -- System clock
    reset    : in  std_logic;        -- Active-high reset
    a        : in  std_logic_vector(WIDTH-1 downto 0); -- Multiplicand
    b        : in  std_logic_vector(WIDTH-1 downto 0); -- Multiplier
    y        : out std_logic_vector(9 downto 0)        -- 10-bit product output
  );
end entity;

architecture Behavioral of pipeline_multiplier is
  -- Pipeline registers for propagating multiplicand
  signal a2_reg, a3_reg, a4_reg      : std_logic_vector(WIDTH-1 downto 0);
  signal a1, a2_next, a3_next, a4_next : std_logic_vector(WIDTH-1 downto 0);

  -- Shifted multiplier bits for partial product generation
  signal b1          : std_logic_vector(4 downto 1); -- Stage1 multiplier slice
  signal b2_next, b2_reg : std_logic_vector(4 downto 2); -- Stage2 multiplier slice
```

```

signal b3_next, b3_reg: std_logic_vector(4 downto 3); -- Stage3 multiplier slice
signal b4_next, b4_reg: std_logic_vector(4 downto 4); -- Stage4 multiplier slice

-- Bit-vector replication and partial products
signal bv0, bv1, bv2, bv3, bv4: std_logic_vector(4 downto 0); -- Bit replication buses
signal bp0, bp1, bp2, bp3, bp4: std_logic_vector(5 downto 0); -- Partial products

-- Pipelined accumulated results
signal pp0          : std_logic_vector(5 downto 0); -- Stage0 partial product
signal pp1_next, pp1_reg: std_logic_vector(6 downto 0); -- Stage1 accumulation
signal pp2_next, pp2_reg: std_logic_vector(7 downto 0); -- Stage2 accumulation
signal pp3_next, pp3_reg: std_logic_vector(8 downto 0); -- Stage3 accumulation
signal pp4_next, pp4_reg: std_logic_vector(9 downto 0); -- Stage4 final result

begin
-- Pipeline Stage 0-1: Initial Partial Product Generation
-- Generate LSB partial product (b[0] * a)
bv0 <= (others => b(0)); -- Replicate b[0] for AND operation
bp0 <= "0" & (bv0 and a); -- Zero-extended AND result
pp0 <= bp0; -- Direct partial product assignment
a1 <= a; -- First stage multiplicand
b1 <= b(4 downto 1); -- Shifted multiplier bits

-- Pipeline Stage 1-2: First Accumulation Stage
-- Generate b[1] partial product and accumulate
bv1 <= (others => b1(1)); -- Replicate b[1] bit
bp1 <= "0" & (bv1 and a1); -- Shifted partial product
pp1_next(6 downto 1) <= ("0" & pp0(5 downto 1)) + bp1; -- Accumulate with shift
pp1_next(0) <= pp0(0); -- Carry LSB through pipeline
a2_next <= a1; -- Propagate multiplicand
b2_next <= b1(4 downto 2); -- Shift multiplier slice

-- Pipeline Stage 2-3: Second Accumulation Stage
bv2 <= (others => b2_reg(2)); -- Replicate current multiplier bit
bp2 <= "0" & (bv2 and a2_reg); -- Generate shifted partial product
pp2_next(7 downto 2) <= ("0" & pp1_reg(6 downto 2)) + bp2; -- Accumulate
pp2_next(1 downto 0) <= pp1_reg(1 downto 0); -- Preserve lower bits
a3_next <= a2_reg; -- Continue multiplicand propagation
b3_next <= b2_reg(4 downto 3); -- Shift multiplier slice

-- Pipeline Stage 3-4: Third Accumulation Stage
bv3 <= (others => b3_reg(3)); -- Current multiplier bit replication
bp3 <= "0" & (bv3 and a3_reg); -- Shifted partial product
pp3_next(8 downto 3) <= ("0" & pp2_reg(7 downto 3)) + bp3; -- Accumulate
pp3_next(2 downto 0) <= pp2_reg(2 downto 0); -- Preserve lower bits
a4_next <= a3_reg; -- Final multiplicand propagation
b4_next <= b3_reg(4 downto 4); -- Final multiplier bit

-- Pipeline Stage 4-5: Final Accumulation Stage
bv4 <= (others => b4_reg(4)); -- MSB replication
bp4 <= "0" & (bv4 and a4_reg); -- MSB partial product
pp4_next(9 downto 4) <= ("0" & pp3_reg(8 downto 4)) + bp4; -- Final sum
pp4_next(3 downto 0) <= pp3_reg(3 downto 0); -- Preserve LSBs

```



```

-- Pipeline Register Update Process
pipe_reg: process(clk, reset)
begin
    if reset = '1' then
        -- Asynchronous reset clears all pipeline registers
        pp1_reg <= (others => '0');
        pp2_reg <= (others => '0');
        pp3_reg <= (others => '0');
        pp4_reg <= (others => '0');
        a2_reg <= (others => '0');
        a3_reg <= (others => '0');
        a4_reg <= (others => '0');
        b2_reg <= (others => '0');
        b3_reg <= (others => '0');
        b4_reg <= (others => '0');
    elsif rising_edge(clk) then
        -- Clock-driven pipeline propagation
        pp1_reg <= pp1_next;
        pp2_reg <= pp2_next;
        pp3_reg <= pp3_next;
        pp4_reg <= pp4_next;
        a2_reg <= a2_next;
        a3_reg <= a3_next;
        a4_reg <= a4_next;
        b2_reg <= b2_next;
        b3_reg <= b3_next;
        b4_reg <= b4_next;
    end if;
end process;

-- Connect final pipeline stage to output
y <= pp4_reg;

end architecture Behavioral;

```

After synthesis and implementations, we can get the schematic diagram as follows:




Figure 5: Pipelined Design Schematic

Using the testbench in the appendix, we can get the result. It computes rightly.




Figure 6: Pipelined Design Result

5 Code Appendix

Testbench codes are as follows, I test $0*0$, $1*1$, $31*31$, $31*1$, $1*31$, $4*8$.

```
-- Initialize and reset
RESET <= '1';
wait for CLK_PERIOD * 2;
RESET <= '0';
wait for CLK_PERIOD;
wait for 100 ns;

-- Test 1: 0 * 0
a_in <= "00000"; b_in <= "00000";
start <= '1';
wait for CLK_PERIOD;
start <= '0';
expected_r := std_logic_vector(to_unsigned(0, 10));
wait until ready = '1';
assert r = expected_r report "Test 1 failed: 0*0" severity error;

-- Test 2: 1 * 1
a_in <= "00001"; b_in <= "00001";
start <= '1';
wait for CLK_PERIOD;
start <= '0';
expected_r := std_logic_vector(to_unsigned(1, 10));
wait until ready = '1';
assert r = expected_r report "Test 2 failed: 1*1" severity error;

-- Test 3: Max value * Max value (31*31)
a_in <= "11111"; b_in <= "11111";
start <= '1';
wait for CLK_PERIOD;
start <= '0';
expected_r := std_logic_vector(to_unsigned(31*31, 10));
wait until ready = '1';
assert r = expected_r report "Test 3 failed: 31*31" severity error;

-- Test 4: Max * 1
a_in <= "11111"; b_in <= "00001";
start <= '1';
wait for CLK_PERIOD;
start <= '0';
expected_r := std_logic_vector(to_unsigned(31, 10));
wait until ready = '1';
assert r = expected_r report "Test 4 failed: 31*1" severity error;

-- Test 5: 1 * Max
a_in <= "00001"; b_in <= "11111";
start <= '1';
wait for CLK_PERIOD;
start <= '0';
expected_r := std_logic_vector(to_unsigned(31, 10));
wait until ready = '1';
assert r = expected_r report "Test 5 failed: 1*31" severity error;
```

```

-- Test 6: Power of 2 multiplication
a_in <= "00100"; b_in <= "01000"; -- 4 * 8
start <= '1';
wait for CLK_PERIOD;
start <= '0';
expected_r := std_logic_vector(to_unsigned(32, 10));
wait until ready = '1';
assert r = expected_r report "Test 6 failed: 4*8" severity error;

-- End of test

```

6 Analysis

6.1 Speed

The multiplier of the **combinatorial design** is based on a purely combinational logic circuit implementation. It generates the output by passing the input signal directly through the combinational logic circuit without the constraints of a clock signal, thus allowing the fastest arithmetic speed in theory. The output signal can be generated almost immediately as long as the input signal is stable and the delay depends mainly on the propagation delay of the combinational logic circuit.

Repetitive-addition design is a multiplication implementation based on addition and shift operations. It accomplishes the multiplication operation by multiple addition operations, each of which corresponds to one bit of the multiplication. Therefore, it is relatively slow and the operation time is proportional to the number of bits of the multiplier.

The **pipeline design** achieves parallel processing by breaking the multiplication operation into multiple stages and inserting registers between each stage. Each stage can be completed in a single clock cycle, so the pipelined design can significantly increase the throughput rate of the multiplier.

Therefore, in terms of operating speed:

$$\text{Combinatorial Design} > \text{Pipeline Design} > \text{Repetitive-Addition Design} \quad (4)$$

Test Case	0*0	1*1	31*31	31*1	1*31	4*8
Combinational	10	20	320	20	320	90
Repetitive-Addition	3.137	4.547	9.162	8.103	8.147	5.693

Table 1: Time of Combinational Design and Repetitive-Addition Design (Unit:ns)

6.2 Space Usage

The **combinatorial design** requires a large number of logic gates to implement multiplication operations, especially for multi-digit numbers. It requires the implementation of the complete multiplication logic, including operations such as partial product generation and summation, and therefore consumes a large amount of hardware resources.

The hardware resource requirements of a **repetitive-addition design** are relatively low because it relies heavily on adders and shift registers. The adder is relatively simple to implement and the shift registers require less logic resources.

A **pipelined design** requires registers to be inserted between each stage and therefore takes up additional hardware resources. In addition, each stage also needs to implement some of the multiplication logic, so the overall hardware resource requirements are higher than for a repeated addition design, but usually lower than for a combinatorial design.

Therefore, in terms of space usage:

$$\text{Combinatorial Design} > \text{Pipeline Design} > \text{Repetitive-Addition Design} \quad (5)$$

6.3 Throughput

The multiplier of a **combinatorial design** can theoretically achieve the fastest arithmetic speed because it is based on a purely combinational logic circuit implementation without the constraint of a clock signal. The output signal can be generated almost immediately as long as the input signal is stable. Therefore, the throughput of a combinational design depends mainly on the propagation delay of the combinational logic circuit.

The multiplier of the **repetitive-addition** design accomplishes the multiplication operation by multiple addition operations, each of which corresponds to one bit of the multiplication. Therefore, its throughput is relatively slow and the operation time is proportional to the number of bits in the multiplier.

The **pipelined design** of the multiplier achieves parallel processing by breaking the complex multiplication operation into multiple simple stages and inserting registers between each stage. Each stage can be completed in a single clock cycle, thus enabling parallel processing. As a result, the throughput of a pipelined design can be significantly increased.

Therefore, in terms of throughput:

$$\text{Pipeline Design} > \text{Combinatorial Design} > \text{Repetitive-Addition Design} \quad (6)$$

7 Conclusion

The comparative analysis of combinational, repetitive-addition, and pipelined multipliers reveals distinct performance characteristics. Combinational design excels in single-operation speed but demands significant hardware resources. Repetitive-addition reduces area usage compared to combinational logic, though at the cost of linearly scaled computation time. Pipelined architectures achieve the highest throughput (1 result per clock cycle after pipeline filling) while maintaining moderate resource overhead. For high-speed real-time applications, combinational or pipelined designs are preferable; resource-constrained systems benefit from repetitive-addition. This study underscores the importance of architectural trade-offs in FPGA-based arithmetic unit design.